

dog_app

October 28, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [2]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/*"))
        dog_files = np.array(glob("/data/dog_images/*/*/*"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

1.2 Explore and Visualize the Datasets

Human images are distributed in 5749 folders named after human names such as "Jonathan Tiomkin", "Charlie Williams", etc. There are 13233 human face images. Images are not evenly distributed among the folders.

The dog dataset is structured into a training set, a validation set, and a test set. Each set comprises 133 breeds of dogs organized into subdirectories by breed name. Later, we will see that the number of images per dog breed is imbalanced. There are 6680 dog images in the training set, 835 images in the validation set, and 836 images in the test set.

```
In [3]: import os
```

```
# Data Exploration
print('Number of folders in /lfw:', len(glob('/data/lfw/*')))
print('Number of human images:', len(glob("/data/lfw/*/*")))
print('Number of folders in /dog_images:', len(glob('/data/dog_images/*')))
print('Folder names in /dog_images:', end=' ')
print(*[x.split('/')[1] for x in glob('/data/dog_images/*')], sep=', ')
print('Number of folders (breed classes) in /dog_images/train:', len(glob("/data/dog_images/train/*")))
print('Number of images in /dog_images/train:', len(glob("/data/dog_images/train/*/*")))
print('Number of folders (breed classes) in /dog_images/valid:', len(glob("/data/dog_images/valid/*")))
print('Number of images in /dog_images/valid:', len(glob("/data/dog_images/valid/*/*")))
print('Number of folders (breed classes) in /dog_images/test:', len(glob("/data/dog_images/test/*")))
print('Number of images in /dog_images/test:', len(glob("/data/dog_images/test/*/*")))
```

```
Number of folders in /lfw: 5749
Number of human images: 13233
Number of folders in /dog_images: 3
Folder names in /dog_images: train, test, valid
Number of folders (breed classes) in /dog_images/train: 133
Number of images in /dog_images/train: 6680
Number of folders (breed classes) in /dog_images/valid: 133
Number of images in /dog_images/valid: 835
Number of folders (breed classes) in /dog_images/test: 133
Number of images in /dog_images/test: 836
```

The dog dataset contains 133 classes. Let print out the first 10 classes.

```
In [4]: classes = sorted([x.split('/')[ -1] for x in glob("/data/dog_images/train/*")])
        print(classes[:10])


['001.Affenpinscher', '002.Afghan_hound', '003.Airedale_terrier', '004.Akita', '005.Alaskan_mala
```

Let's look at a couple of images from the human and the dog datasets.

```
In [5]: from PIL import Image
        import matplotlib.pyplot as plt
        %matplotlib inline

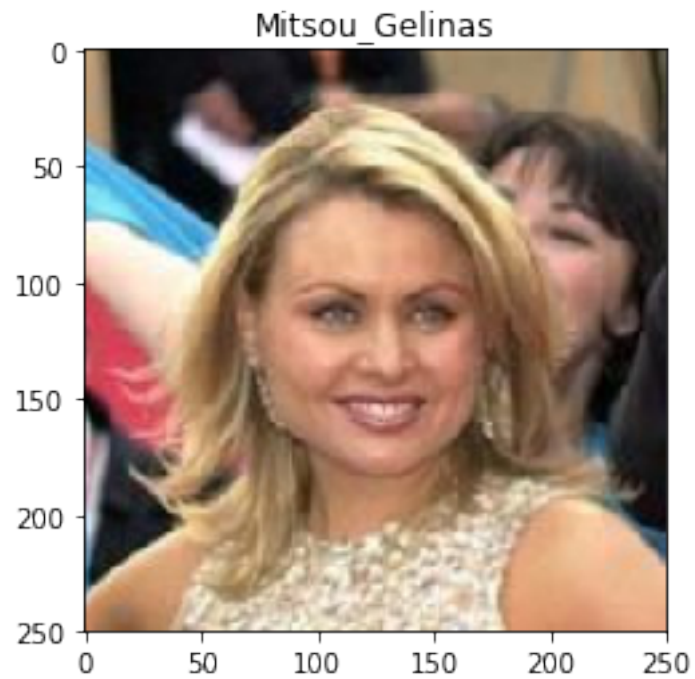
        img = Image.open(human_files[0])
        title = human_files[0].split('/')[ -2]
        plt.title(title)
        plt.imshow(img)

Out[5]: <matplotlib.image.AxesImage at 0x7f1122168400>
```




```
In [6]: img = Image.open(human_files[100])  
        title = human_files[100].split('/')[-2]  
        plt.title(title)  
        plt.imshow(img)
```

```
Out[6]: <matplotlib.image.AxesImage at 0x7f112210aa58>
```



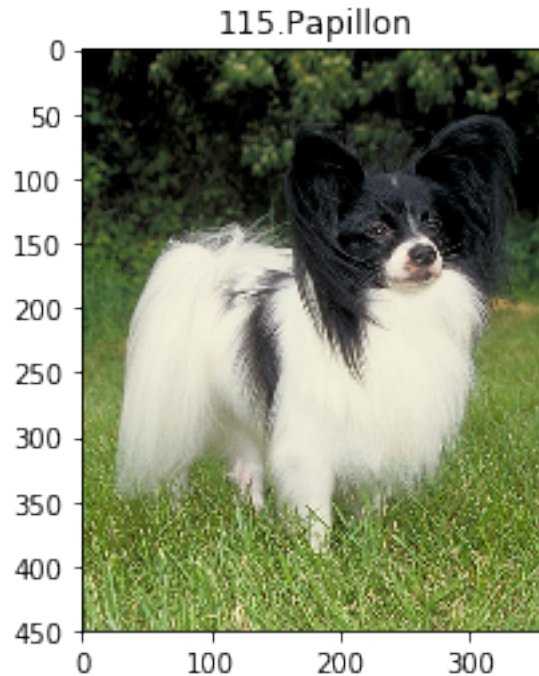
```
In [7]: img = Image.open(dog_files[0])  
        title = dog_files[0].split('/')[-2]  
        plt.title(title)  
        plt.imshow(img)
```

```
Out[7]: <matplotlib.image.AxesImage at 0x7f112086b7f0>
```



```
In [8]: img = Image.open(dog_files[500])
        title = dog_files[500].split('/')[-2]
        plt.title(title)
        plt.imshow(img)
```

```
Out[8]: <matplotlib.image.AxesImage at 0x7f11207d6208>
```



The dog dataset is imbalanced. This observation is critical to use other metrics instead of accuracy that distinguishes correctly between the numbers of correctly classified examples of different classes i.e., precision and recall.

```
In [9]: num_images_per_folder = {}
        num_images_per_folder['train'] = [len(glob(x+'/*')) for x in glob("/data/dog_images/train")]
        num_images_per_folder['valid'] = [len(glob(x+'/*')) for x in glob("/data/dog_images/valid")]
        num_images_per_folder['test'] = [len(glob(x+'/*')) for x in glob("/data/dog_images/test")]
        num_images_avg = {}
        num_images_avg['train'] = sum(num_images_per_folder['train'])/len(num_images_per_folder['train'])
        num_images_avg['valid'] = sum(num_images_per_folder['valid'])/len(num_images_per_folder['valid'])
        num_images_avg['test'] = sum(num_images_per_folder['test'])/len(num_images_per_folder['test'])
```


```
In [10]: print('Average number of images per dog breed in training set is', num_images_avg['train'])
         print('Average number of images per dog breed in validation set is', num_images_avg['valid'])
         print('Average number of images per dog breed in test set is', num_images_avg['test'])
```

```
Average number of images per dog breed in training set is 50.225563909774436
Average number of images per dog breed in validation set is 6.2781954887218046
Average number of images per dog breed in test set is 6.285714285714286
```

```
In [11]: import matplotlib.pyplot as plt
         %matplotlib inline

         plt.figure(figsize=(10, 30))
```

```
bar1 = plt.barh(classes, num_images_per_folder['train'])
#plt.barh(classes, num_images_per_folder['test'])
bar2 = plt.barh(classes, num_images_per_folder['valid'])
line1 = plt.axvline(num_images_avg['train'], color='blue')
line2 = plt.axvline(num_images_avg['valid'], color='orange')
plt.ylabel('Dog breed classes')
plt.xlabel('Number of images per folder')
plt.title('Dog breed class distributions')
plt.legend((bar1, bar2, line1, line2), ('train', 'valid', 'avg train', 'avg valid'))
plt.show()
```

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [12]: import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# extract pre-trained face detector
face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

# load color (BGR) image
img = cv2.imread(human_files[0])
# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)


# print number of faces detected in the image
print('Number of faces detected:', len(faces))

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()
```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.2.1 Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [13]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.2.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

98 of the first 100 images in `human_files` or 98% have a detected human face.

17 of the first 100 images in `dog_files` or 17% have a detected human face.

```
In [14]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]
```

```
##-## Do NOT modify the code above this line. ##-##
```

```
## TODO: Test the performance of the face_detector algorithm
```

```
## on the images in human_files_short and dog_files_short.
```

```
human_count = sum([face_detector(human) for human in human_files_short])
```

```
dog_count = sum([face_detector(dog) for dog in dog_files_short])
```

```
print('{} of the first 100 images in human_files have a detected human face.'.format(hu_
```

```
print('{} of the first 100 images in dog_files have a detected human face.'.format(dog_
```

98 of the first 100 images in `human_files` have a detected human face.

17 of the first 100 images in `dog_files` have a detected human face.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [15]: ### (Optional)
```

```
### TODO: Test performance of another face detection algorithm.
```

```
### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.2.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [16]: import torch
import torchvision.models as models

# define VGG16 model
VGG16 = models.vgg16(pretrained=True)

# check if CUDA is available
use_cuda = torch.cuda.is_available()

# move model to GPU if CUDA is available
if use_cuda:
    VGG16 = VGG16.cuda()
```

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth [100%| 553433881/553433881 [00:05<00:00, 107552511.85it/s]

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.2.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [17]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image
    transform = transforms.Compose([
        transforms.Resize(256),
```

```

        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    img = Image.open(img_path).convert('RGB')
    img = transform(img)
    img = img.unsqueeze(0)

    if use_cuda:
        img = img.cuda()

    VGG16.eval()
    output = VGG16(img)

    return torch.max(output,1)[1].item() # predicted class index

```

1.2.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```

In [18]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    index = VGG16_predict(img_path)
    return index >= 151 and index <= 268 # true/false

```

1.2.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

Percentage of detected dog on `human_files_short` is 0.0%

Percentage of detected dog on `dog_files_short` is 100.0%

```

In [19]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
human_count = 0
dog_count = 0

for i in human_files_short:
    if dog_detector(i):

```

```

        human_count+=1
    for i in dog_files_short:
        if dog_detector(i):
            dog_count+=1

    print("Percentage of detected dog on human_files_short is {}".format(100*human_count/100.0))
    print("Percentage of detected dog on dog_files_short is {}".format(100*dog_count/100.0))

```

Percentage of detected dog on human_files_short is 0.0%
 Percentage of detected dog on dog_files_short is 100.0%

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on human_files_short and dog_files_short.

```

In [20]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.

```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.2.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [21]: import os
         from torchvision import datasets

         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         # Define batch size
         batch_size = 8
         num_workers = 0

         # directories
         data_dir = '/data/dog_images/'
         train_dir = os.path.join(data_dir, 'train/')
         valid_dir = os.path.join(data_dir, 'valid/')
         test_dir = os.path.join(data_dir, 'test/')

         # pre-process data
         data_transforms = {'train': transforms.Compose([transforms.Resize(256),
                                                         transforms.RandomHorizontalFlip(),
                                                         transforms.RandomRotation(10),
                                                         transforms.CenterCrop(224),
                                                         transforms.ToTensor(),
                                                         transforms.Normalize((0.5, 0.5, 0.5), (
                                                         'valid': transforms.Compose([transforms.Resize(256),
                                                         transforms.CenterCrop(224),
                                                         transforms.ToTensor(),
```



```

transformations.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))

        'test': transformations.Compose([transformations.Resize(256),
                                          transformations.CenterCrop(224),
                                          transformations.ToTensor(),
                                          transformations.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
                                          ])
    }

    train_data = datasets.ImageFolder(train_dir, transform=data_transforms['train'])
    valid_data = datasets.ImageFolder(valid_dir, transform=data_transforms['valid'])
    test_data = datasets.ImageFolder(test_dir, transform=data_transforms['test'])

    train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_workers=4)
    valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_workers=4)
    test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers=4)

    loaders_scratch = {'train': train_loader, 'valid': valid_loader, 'test': test_loader}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

In the original [VGG-16 paper](#), the authors chose fixed-size 224x224 images as the input. Thus I rescaled the original images to 256x256 before cropping to get 224x224 images. The image rescaling is critical since cropping a 224x224 image out of a much larger original is unlikely to contain the features we are interested in to train the model.

Data augmentation is a way to extend a dataset and improve generalization to avoid overfitting, therefore, I augmented the dataset using random horizontal flipping and random rotation up to 10 degrees. The data augmentation is performed on training set only.

1.2.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [22]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        # convolutional layer (224x224x3 image tensor)
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3, stride=2, padding=1)
        # convolutional layer (56x56x16 image tensor)
        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=2, padding=1)

```

```

        # convolutional layer (14x14x32 image tensor)
        self.conv3 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        # max pooling layer
        self.pool = nn.MaxPool2d(2, 2)
        # linear layer (64 * 7 * 7 -> 512)
        self.fc1 = nn.Linear(64 * 7 * 7, 512)
        # linear layer (512 -> 512)
        self.fc2 = nn.Linear(512, 512)
        # linear layer (512 -> 133)
        self.fc3 = nn.Linear(512, 133)
        # dropout layer (p=0.25)
        self.dropout = nn.Dropout(0.25)

    def forward(self, x):
        ## Define forward behavior
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        # flatten image input
        x = x.view(-1, 64 * 7 * 7)
        # add dropout layer
        x = self.dropout(x)
        # add 1st hidden layer, with relu activation function
        x = F.relu(self.fc1(x))
        # add dropout layer
        x = self.dropout(x)
        # add 2nd hidden layer, with relu activation function
        x = self.fc2(x)
        # add dropout layer
        x = self.dropout(x)
        # add 3rd hidden layer, with relu activation function
        x = self.fc3(x)

        return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))

```

```

(conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(fc1): Linear(in_features=3136, out_features=512, bias=True)
(fc2): Linear(in_features=512, out_features=512, bias=True)
(fc3): Linear(in_features=512, out_features=133, bias=True)
(dropout): Dropout(p=0.25)
)

```

```

In [23]: !pip install torchsummary
         from torchsummary import summary
         summary(model_scratch, (3, 224, 224))

```

Collecting torchsummary

Downloading <https://files.pythonhosted.org/packages/7d/18/1474d06f721b86e6a9b9d7392ad68bed711a>

Installing collected packages: torchsummary

Successfully installed torchsummary-1.5.1

```

-----
Layer (type)              Output Shape              Param #
=====
      Conv2d-1             [-1, 16, 112, 112]         448
    MaxPool2d-2            [-1, 16, 56, 56]           0
      Conv2d-3             [-1, 32, 28, 28]        4,640
    MaxPool2d-4            [-1, 32, 14, 14]           0
      Conv2d-5             [-1, 64, 14, 14]       18,496
    MaxPool2d-6            [-1, 64, 7, 7]            0
      Dropout-7            [-1, 3136]                 0
      Linear-8              [-1, 512]        1,606,144
      Dropout-9            [-1, 512]                 0
      Linear-10             [-1, 512]        262,656
      Dropout-11           [-1, 512]                 0
      Linear-12             [-1, 133]         68,229
=====

```

Total params: 1,960,613

Trainable params: 1,960,613

Non-trainable params: 0

```

-----
Input size (MB): 0.57
Forward/backward pass size (MB): 2.31
Params size (MB): 7.48
Estimated Total Size (MB): 10.37
-----

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

The scratch model is structured as follows:

3 convolutional (conv) layers transform a 3-channel image to a 64-channel feature map. All conv layers have a kernel size of 3 and a padding of 1 (to preserve spatial resolution after convolution). The first 2 conv layers have stride 2 while the third one has stride 1.

Max-pooling is performed over a 2 x 2-pixel window, with stride 2 after conv layers which halves the height and width. The feature map gets smaller as we add more layers until we are finally left with a small feature map, which can be flattened into a vector.

3 fully connected layers at the end produce an output vector of size 133 for each image.

3 dropout layers of 0.25 are added to avoid overfitting.

All hidden layers are equipped with the rectification (ReLU) non-linearity.

1.2.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [24]: import torch.optim as optim
```

```
### TODO: select loss function
criterion_scratch = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.02)
# optimizer_scratch = optim.Adam(model_scratch.parameters(), lr = 0.01)
```

1.2.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```
In [14]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
```

```

    ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    # clear the gradients of all optimized variables
    optimizer.zero_grad()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # backward pass: compute gradient of the loss with respect to model parameters
    loss.backward()
    # perform a single optimization step (parameter update)
    optimizer.step()
    # update average training loss
    train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    # forward pass: compute predicted outputs by passing inputs to the model
    with torch.no_grad():
        output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # update average validation loss
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
        train_loss, valid_loss))
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model

```

```

# train the model
model_scratch = train(15, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.885879      Validation Loss: 4.874210
Validation loss decreased (inf --> 4.874210). Saving model ...
Epoch: 2      Training Loss: 4.866884      Validation Loss: 4.825825
Validation loss decreased (4.874210 --> 4.825825). Saving model ...
Epoch: 3      Training Loss: 4.694588      Validation Loss: 4.615354
Validation loss decreased (4.825825 --> 4.615354). Saving model ...
Epoch: 4      Training Loss: 4.531321      Validation Loss: 4.465499
Validation loss decreased (4.615354 --> 4.465499). Saving model ...
Epoch: 5      Training Loss: 4.371618      Validation Loss: 4.374201
Validation loss decreased (4.465499 --> 4.374201). Saving model ...
Epoch: 6      Training Loss: 4.265738      Validation Loss: 4.324693
Validation loss decreased (4.374201 --> 4.324693). Saving model ...
Epoch: 7      Training Loss: 4.158158      Validation Loss: 4.183197
Validation loss decreased (4.324693 --> 4.183197). Saving model ...
Epoch: 8      Training Loss: 4.071888      Validation Loss: 4.128459
Validation loss decreased (4.183197 --> 4.128459). Saving model ...
Epoch: 9      Training Loss: 3.981011      Validation Loss: 4.072647
Validation loss decreased (4.128459 --> 4.072647). Saving model ...
Epoch: 10     Training Loss: 3.893157      Validation Loss: 4.061979
Validation loss decreased (4.072647 --> 4.061979). Saving model ...
Epoch: 11     Training Loss: 3.785850      Validation Loss: 3.952812
Validation loss decreased (4.061979 --> 3.952812). Saving model ...
Epoch: 12     Training Loss: 3.697476      Validation Loss: 3.902146
Validation loss decreased (3.952812 --> 3.902146). Saving model ...
Epoch: 13     Training Loss: 3.559261      Validation Loss: 3.842603
Validation loss decreased (3.902146 --> 3.842603). Saving model ...
Epoch: 14     Training Loss: 3.480124      Validation Loss: 3.890282
Epoch: 15     Training Loss: 3.361239      Validation Loss: 3.780858
Validation loss decreased (3.842603 --> 3.780858). Saving model ...

```

1.2.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [25]: def test(loaders, model, criterion, use_cuda):
```

```

    # monitor test loss and accuracy
    test_loss = 0.

```

```

correct = 0.
total = 0.

predicted_labels = []
true_labels = []

model.eval()
for batch_idx, (data, target) in enumerate(loaders['test']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    predicted_labels.append(pred)
    true_labels.append(target)
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

return predicted_labels, true_labels

```

In [27]: # call test function

```
predicted_labels, true_labels = test(loaders_scratch, model_scratch, criterion_scratch,
```

Test Loss: 3.773037

Test Accuracy: 13% (115/836)

In [28]: predictions = []

```

for labels in predicted_labels:
    for label in np.array(labels):
        predictions.append(label[0])

```

```

ground_truths = []
for labels in true_labels:

```

```

        for label in np.array(labels):
            ground_truths.append(label)

In [29]: from sklearn.metrics import confusion_matrix, precision_recall_fscore_support
precision = np.mean(precision_recall_fscore_support(ground_truths, predictions)[0])
recall = np.mean(precision_recall_fscore_support(ground_truths, predictions)[1])
print('Average test precision: {:.1f}%'.format(precision*100))
print('Average test recall: {:.1f}%'.format(recall*100))

Average test precision: 12.2%
Average test recall: 12.3%

/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning
'precision', 'predicted', average, warn_for)

```

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)
 You will now use transfer learning to create a CNN that can identify dog breed from images.
 Your CNN must attain at least 60% accuracy on the test set.

1.2.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [30]: ## TODO: Specify data loaders
        batch_size = 32

        train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_workers=4)
        valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size, num_workers=4)
        test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers=4)

        loaders_transfer = {'train': train_loader, 'valid': valid_loader, 'test': test_loader}

```

1.2.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [31]: import torchvision.models as models
        import torch.nn as nn

        ## TODO: Specify model architecture

```



```

# Load the pretrained model from PyTorch
model_transfer = models.vgg16(pretrained=True)

# Freeze training for all "features" layers
for param in model_transfer.features.parameters():
    param.requires_grad = False

n_inputs = model_transfer.classifier[6].in_features

# add last linear layer (n_inputs -> 133 dog breed classes)
# new layers automatically have requires_grad = True
model_transfer.classifier[6] = nn.Linear(n_inputs, 133)

# print out the model structure

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

The CNN model built from scratch can be improved significantly by transfer learning. The idea is to reuse the lower layers of a pre-trained model for feature extraction and only train a few higher layers for classification. Modern CNNs training on huge datasets like ImageNet have great performance. Here I am interested in a couple of state-of-the-art pre-trained architectures VGGNet (Zisserman, 2015) and (K. He, 2016) as a fixed feature extractor for my models.

VGGNet is worth trying because it has a very simple and classical architecture and has great performance, coming in second in the ImageNet 2014 competition. VGGNet is already adopted in the dog detection task and thus can be reused. The idea here is that I keep all the convolutional layers, but replace the final fully-connected layer with my classifier. This way I can use VGGNet as a fixed feature extractor for the images then easily train a simple classifier on top of that.

1.2.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

In [32]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)

```

1.2.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [19]: n_epochs = 20

# train the model
model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,

```

```
# load the model that got the best validation accuracy (uncomment the line below)  
model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1      Training Loss: 4.433197      Validation Loss: 3.660359  
Validation loss decreased (inf --> 3.660359). Saving model ...  
Epoch: 2      Training Loss: 3.068509      Validation Loss: 2.253841  
Validation loss decreased (3.660359 --> 2.253841). Saving model ...  
Epoch: 3      Training Loss: 1.973345      Validation Loss: 1.419213  
Validation loss decreased (2.253841 --> 1.419213). Saving model ...  
Epoch: 4      Training Loss: 1.433813      Validation Loss: 1.035404  
Validation loss decreased (1.419213 --> 1.035404). Saving model ...  
Epoch: 5      Training Loss: 1.149600      Validation Loss: 0.854170  
Validation loss decreased (1.035404 --> 0.854170). Saving model ...  
Epoch: 6      Training Loss: 1.004561      Validation Loss: 0.736856  
Validation loss decreased (0.854170 --> 0.736856). Saving model ...  
Epoch: 7      Training Loss: 0.915746      Validation Loss: 0.695458  
Validation loss decreased (0.736856 --> 0.695458). Saving model ...  
Epoch: 8      Training Loss: 0.825299      Validation Loss: 0.636218  
Validation loss decreased (0.695458 --> 0.636218). Saving model ...  
Epoch: 9      Training Loss: 0.768361      Validation Loss: 0.605259  
Validation loss decreased (0.636218 --> 0.605259). Saving model ...  
Epoch: 10     Training Loss: 0.729283      Validation Loss: 0.576355  
Validation loss decreased (0.605259 --> 0.576355). Saving model ...  
Epoch: 11     Training Loss: 0.680769      Validation Loss: 0.567901  
Validation loss decreased (0.576355 --> 0.567901). Saving model ...  
Epoch: 12     Training Loss: 0.657427      Validation Loss: 0.544684  
Validation loss decreased (0.567901 --> 0.544684). Saving model ...  
Epoch: 13     Training Loss: 0.618902      Validation Loss: 0.539705  
Validation loss decreased (0.544684 --> 0.539705). Saving model ...  
Epoch: 14     Training Loss: 0.591450      Validation Loss: 0.519682  
Validation loss decreased (0.539705 --> 0.519682). Saving model ...  
Epoch: 15     Training Loss: 0.577838      Validation Loss: 0.509746  
Validation loss decreased (0.519682 --> 0.509746). Saving model ...  
Epoch: 16     Training Loss: 0.557196      Validation Loss: 0.505338  
Validation loss decreased (0.509746 --> 0.505338). Saving model ...  
Epoch: 17     Training Loss: 0.528556      Validation Loss: 0.495418  
Validation loss decreased (0.505338 --> 0.495418). Saving model ...  
Epoch: 18     Training Loss: 0.504658      Validation Loss: 0.490786  
Validation loss decreased (0.495418 --> 0.490786). Saving model ...  
Epoch: 19     Training Loss: 0.493439      Validation Loss: 0.484817  
Validation loss decreased (0.490786 --> 0.484817). Saving model ...  
Epoch: 20     Training Loss: 0.494086      Validation Loss: 0.481286  
Validation loss decreased (0.484817 --> 0.481286). Saving model ...
```

1.2.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [34]: predicted_labels, true_labels = test(loaders_transfer, model_transfer, criterion_transfer)
```

Test Loss: 0.528622

Test Accuracy: 83% (701/836)

```
In [35]: predictions = []
        for labels in predicted_labels:
            for label in np.array(labels):
                predictions.append(label[0])

        ground_truths = []
        for labels in true_labels:
            for label in np.array(labels):
                ground_truths.append(label)

        precision = np.mean(precision_recall_fscore_support(ground_truths, predictions)[0])
        recall = np.mean(precision_recall_fscore_support(ground_truths, predictions)[1])
        print('Average test precision: {:.1f}%'.format(precision*100))
        print('Average test recall: {:.1f}%'.format(recall*100))
```

Average test precision: 85.0%

Average test recall: 82.9%

On the other hand, ResNet is the winner of the ImageNet 2015 competition and has amazing performance on image classification. It helps to train very deep networks by using skip connections (or shortcut connections). Specifically, it adds the original input back to the output feature map obtained by passing the input through one or more conv layers. This seemingly small change produces a drastic improvement in the performance of the model.

We will use the 101-layer-deep ResNet architecture (ResNet101) pre-trained on the ImageNet dataset. Similar to VGGNet, we only need to replace the final fully-connected layer with our classifier to produce 133-dimensional output. The following code block implements the solution.

```
In [21]: # import torch.optim as optim

        model_transfer2 = models.resnet101(pretrained=True)

        for param in model_transfer2.parameters():
            param.requires_grad = False

        model_transfer2.fc = nn.Linear(2048, 133)
```

```

if use_cuda:
    model_transfer2 = model_transfer2.cuda()

criterion_transfer2 = nn.CrossEntropyLoss()
optimizer_transfer2 = optim.Adam(model_transfer2.fc.parameters(), lr=0.001)

# train the model
model_transfer2 = train(20, loaders_transfer, model_transfer2, optimizer_transfer2, cri

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer2.load_state_dict(torch.load('model_transfer2.pt'))

```

Downloading: "https://download.pytorch.org/models/resnet101-5d3b4d8f.pth" to /root/.torch/models
100%|| 178728960/178728960 [00:14<00:00, 12669644.24it/s]

Epoch: 1	Training Loss: 1.922533	Validation Loss: 0.728569
Validation loss decreased (inf --> 0.728569). Saving model ...		
Epoch: 2	Training Loss: 0.610473	Validation Loss: 0.585945
Validation loss decreased (0.728569 --> 0.585945). Saving model ...		
Epoch: 3	Training Loss: 0.439258	Validation Loss: 0.513522
Validation loss decreased (0.585945 --> 0.513522). Saving model ...		
Epoch: 4	Training Loss: 0.350083	Validation Loss: 0.462946
Validation loss decreased (0.513522 --> 0.462946). Saving model ...		
Epoch: 5	Training Loss: 0.291649	Validation Loss: 0.450441
Validation loss decreased (0.462946 --> 0.450441). Saving model ...		
Epoch: 6	Training Loss: 0.270571	Validation Loss: 0.429810
Validation loss decreased (0.450441 --> 0.429810). Saving model ...		
Epoch: 7	Training Loss: 0.232485	Validation Loss: 0.421586
Validation loss decreased (0.429810 --> 0.421586). Saving model ...		
Epoch: 8	Training Loss: 0.217263	Validation Loss: 0.469459
Epoch: 9	Training Loss: 0.193880	Validation Loss: 0.428393
Epoch: 10	Training Loss: 0.190744	Validation Loss: 0.437733
Epoch: 11	Training Loss: 0.187236	Validation Loss: 0.397946
Validation loss decreased (0.421586 --> 0.397946). Saving model ...		
Epoch: 12	Training Loss: 0.158173	Validation Loss: 0.403430
Epoch: 13	Training Loss: 0.155980	Validation Loss: 0.436753
Epoch: 14	Training Loss: 0.149405	Validation Loss: 0.429206
Epoch: 15	Training Loss: 0.146173	Validation Loss: 0.479776
Epoch: 16	Training Loss: 0.139664	Validation Loss: 0.483285
Epoch: 17	Training Loss: 0.125960	Validation Loss: 0.443054
Epoch: 18	Training Loss: 0.129630	Validation Loss: 0.477187
Epoch: 19	Training Loss: 0.137159	Validation Loss: 0.521238
Epoch: 20	Training Loss: 0.123384	Validation Loss: 0.568297

```

In [38]: model_transfer2.load_state_dict(torch.load('model_transfer2.pt'))
         predicted_labels, true_labels = test(loaders_transfer, model_transfer2, criterion_trans

```

Test Loss: 0.521041

Test Accuracy: 85% (718/836)

```
In [39]: predictions = []
         for labels in predicted_labels:
             for label in np.array(labels):
                 predictions.append(label[0])

         ground_truths = []
         for labels in true_labels:
             for label in np.array(labels):
                 ground_truths.append(label)

         precision = np.mean(precision_recall_fscore_support(ground_truths, predictions)[0])
         recall = np.mean(precision_recall_fscore_support(ground_truths, predictions)[1])
         print('Average test precision: {:.1f}%'.format(precision*100))
         print('Average test recall: {:.1f}%'.format(recall*100))
```

Average test precision: 87.1%

Average test recall: 84.3%

```
/opt/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning:
'precision', 'predicted', average, warn_for)
```

1.2.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [40]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

         # list of class names by index, i.e. a name can be accessed like class_names[0]
         class_names = [item[4:].replace("_", " ") for item in train_loader.dataset.classes]

         def load_transform_image(img_path):

             # Pre-process
             transform = transforms.Compose([
                 transforms.Resize(256),
                 transforms.CenterCrop(224),
                 transforms.ToTensor(),
                 transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
             ])
```

```

        # Open the image and apply the transformation process
        img = Image.open(img_path).convert('RGB')
        img = transform(img)[:3,:,:].unsqueeze(0)

    return img

def predict_breed_transfer(img_path, class_names, model):
    # load the image and return the predicted breed
    img = load_transform_image(img_path)

    if use_cuda:
        img = img.cuda()
        model = model.cuda()

    model.eval()
    prediction = model(img)

    index = torch.max(prediction, 1)[1].item()

    return class_names[index]

In [41]: predict_breed_transfer('./images/sample_cnn.png', class_names, model_transfer2)

Out[41]: 'Cavalier king charles spaniel'

In [42]: # Test
         predict_breed_transfer('./images/Welsh_springer_spaniel_08203.jpg', class_names, model_

Out[42]: 'Welsh springer spaniel'

```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.


Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.2.18 (IMPLEMENTATION) Write your Algorithm

```

In [43]: ### TODO: Write your algorithm.
         ### Feel free to use as many code cells as needed.

```



Sample Human Output

```
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    img = Image.open(img_path)
    plt.imshow(img)

    if dog_detector(img_path) is True:
        prediction = predict_breed_transfer(img_path, class_names, model_transfer2)
        print("\n\nHello, Dog!")
        plt.show()
        print("You look like a ...\n{0}".format(prediction))

    elif face_detector(img_path) is True:
        prediction = predict_breed_transfer(img_path, class_names, model_transfer2)
        print("\n\nHello, Human!")
        plt.show()
        print("You look like a ...\n{0}".format(prediction))
    else:
        print("\n\nError! No Dog or Human detected!")
        plt.show()
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.2.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement)

Overall, the dog app works rather well. It can detect human, dog, or neither in the given images. Then it provides estimates on the breed for the dog image, the resembling dog breed for

the human image, or returns an error message.

There is a lot of room to improve our model performance and reduce training time:

Try to increase the training data by other data augmentation techniques like random shift, random size, generated images with various contrasts, etc.

Try other variants of ResNet or different CNN architectures (GoogLeNet, Xception, SEnet – winner of ImageNet 2017, etc.) or a combination of CNN models.


Playing with hyperparameters like batch size, learning rate, optimizer, number of epochs, activation function, etc.

Instead of using a fixed learning rate, try to use a learning rate scheduler, which will change the learning rate after every batch of training.

```
In [44]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.


         ## suggested code, below
         for file in np.hstack((human_files[:3], dog_files[:3])):
             run_app(file)
```

Hello, Human!




You look like a ...
Welsh springer spaniel

Hello, Human!




You look like a ...
Lowchen

Hello, Human!




You look like a ...
Welsh springer spaniel

Hello, Dog!




You look like a ...
Mastiff

Hello, Dog!



You look like a ...
Mastiff

Hello, Dog!



You look like a ...
Bullmastiff

In []: